

AI Aerial Surveillance & Crime Detection System

Full Technical Project Report

September 01, 2026

Flask Web Application | Three-Model AI Pipeline

1. Project Overview

The AI Aerial Surveillance & Crime Detection System is a deep learning platform that automatically analyzes aerial drone footage and detects criminal activity in real time. It is delivered as a Flask web application integrating three specialized AI models in a coordinated two-stage pipeline. The system classifies seven scene categories and produces annotated video output with bounding boxes, confidence scores, and incident reports.

Crime Categories Detected

Class	Description
Normal	No criminal activity — ordinary movement or gathering
Violence	Physical assault, fighting, or armed attack with a blade
Robbery	Snatch-and-run, mugging, or armed theft
Shooting	Active firearm threat — REQUIRES confirmed firearm detection
FireExplosion	Fire, explosion, or explosive device present
Accident	Vehicle collision or unintentional injury event
Vandalism	Property damage without direct person-to-person violence

2. System Architecture

The system operates as a two-stage pipeline. Stage 1 detects people and weapons per frame. Stage 2 classifies crime activity over 8-frame temporal windows. A multi-modal decision engine fuses both stages before producing the final verdict.

	USER UPLOADS VIDEO via Flask Web Interface	
	STAGE 1 — Spatial Detection (YOLOv8) Person Detector Weapon Detector COCO Object Detector	
	STAGE 2 — Temporal Classification (R(2+1)D) 8-frame sliding window crime_aerial_augmented_best1.pth	
	MULTI-MODAL DECISION ENGINE fuse_multimodal_crime_decision() Physical Evidence Gating	
	ANNOTATED VIDEO + JSON REPORT + EVIDENCE CLIPS	

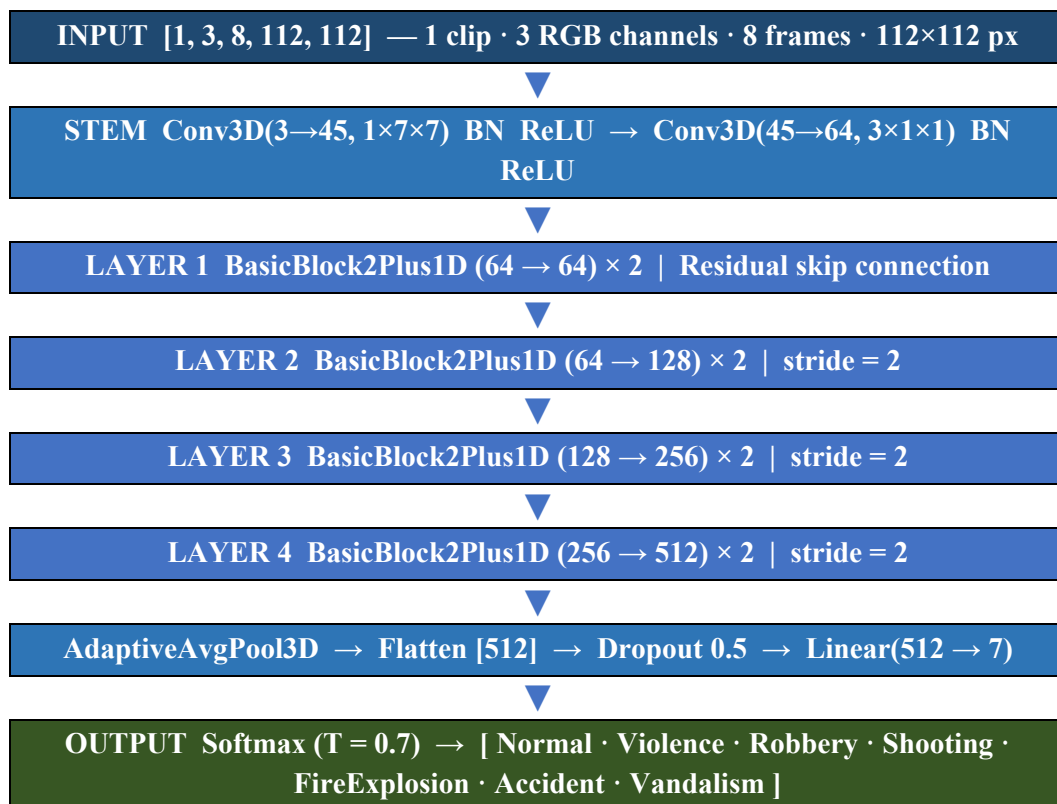
3. Three AI Models

Model	File	Size	Architecture	Output
Crime Classifier	crime_aerial_augmented_best1.pth	358 MB	CrimeR2Plus1D (3D CNN)	7 crime classes
Person Detector	drone_person_detector_best.pt	18.3 MB	YOLOv8 custom aerial	Bounding boxes + track IDs
Weapon Detector	weapon_detector_best.pt	18.3 MB	YOLOv8 specialized	Firearm / Blade / Explosive

3.1 Crime Classifier: crime_aerial_augmented_best1.pth

Primary crime classification model — a 3D Convolutional Neural Network (CrimeR2Plus1D) based on the R(2+1)D architecture, fine-tuned on aerial drone footage. It processes 8-frame video clips and outputs probabilities for 7 crime categories.


CrimeR2Plus1D Architecture



Parameter	Value
Input shape	[1, 3, 8, 112, 112]
State-dict tensors	224
Classes	7
Temperature scaling	T = 0.7
Dropout	0.5 before classifier
Crime threshold	≥ 0.30 and must exceed Normal probability
Video-level crime rule	$\geq 15\%$ windows as crime OR top crime ≥ 0.40
Clarity filter trigger	Confidence < 0.48 $\rightarrow$ aerial CLAHE re-evaluation

3.2 Person Detector: drone_person_detector_best.pt

Custom YOLOv8 fine-tuned on aerial drone footage to detect people from top-down and oblique viewpoints. Works from 5 m to 100 m altitude. Integrated with ByteTrack for persistent multi-object tracking and speed estimation.



Parameter	Value
Confidence threshold	0.20 (permissive for small aerial targets)
NMS IoU threshold	0.40
Walking speed	≤ 8.0 px/frame
Running speed	> 8.0 px/frame
Interaction proximity	Center distance < 75% of max person bbox size

3.3 Weapon Detector: weapon_detector_best.pt

Specialized YOLOv8 trained to detect three weapon categories in aerial footage. Serves as the physical evidence gate — Shooting can only be confirmed if a Firearm is simultaneously detected.

Class	Label	Keywords Matched	Effect on Decision
0	Firearm	gun, pistol, rifle, shotgun, firearm	REQUIRED to confirm Shooting; else forced Normal
1	Blade	knife, blade, dagger, sword, machete	Routes to Violence; or Robbery if person is running
2	Explosive	bomb, grenade, dynamite, explosive	Routes to FireExplosion classification


4. Flask Web Application

app.py (262 lines) is the HTTP layer. It handles video uploads, coordinates the AI pipeline, serves annotated outputs, and renders three Jinja2 template pages.

4.1 Routes

Route	Method	Purpose
/ or /index	GET	Dashboard — upload form and analysis results
/upload	POST	Receive video file, run AI pipeline, return results
/live or /detection	GET	Real-time MJPEG detection stream page
/live_video/<filename>	GET	MJPEG frame stream from the AI engine
/reports	GET	Historical audit reports listing
/outputs/<path>	GET	Serve annotated videos and JSON reports
/evidence/<path>	GET	Serve extracted evidence clips
/download_report/<fn>	GET	Download JSON report as file attachment

4.2 Processing Flow




5. Image Preprocessing (AdaptiveDronePreprocessor)

Mode	Use Case	Algorithms
auto	Automatic — diagnoses scene, selects best mode	Analyzes brightness, contrast, sharpness, colorfulness
aerial_drone	Hazy aerial footage	CLAHE, dark-channel dehaze, unsharp mask, histogram stretch
low_light_night	Night / low light	AGCWD adaptive gamma, shadow lift, noise reduction
yolo_enhanced	Maximize YOLO recall	Edge enhancement, contrast boost, Gaussian sharpen
none	No processing needed	Raw frame passthrough

6. Adaptive Acquisition

When the drone flies at high altitude, people appear very small ($< 3.5\%$ of frame height) and classification confidence drops below 0.55. The module applies multi-scale zoom cropping ($1.5\times - 3.0\times$) and FSRCNN $3\times$ super-resolution to improve accuracy.



7. Output and Reporting

Output Type	Description
Annotated Video	H.264 MP4 with HUD: crime label, confidence bar, person/weapon/vehicle bounding boxes, speed labels, decision reason tag
JSON Report	Video metadata, final verdict, all window-level predictions, entity summary, evidence clip paths, processing metrics
Evidence Clips	MP4 segments of highest-confidence crime windows — saved to outputs/evidence/ with class name and timestamp in filename

8. Technology Stack

Library / Tool	Version	Role
Flask	≥ 3.0.0	HTTP routing, Jinja2 templates, file serving
PyTorch	≥ 2.0.0	Model inference, tensor ops, MPS/CUDA
TorchVision	≥ 0.15.0	ResNet-50 backbone support
Ultralytics YOLOv8	≥ 8.0.0	Person/weapon detection + ByteTrack
OpenCV	≥ 4.8.0	Frame processing, HUD rendering, video I/O
NumPy	≥ 1.24.0	Array operations, normalization
python-docx	≥ 1.2.0	DOCX report generation
FFmpeg (external)	Latest	H.264 video re-encoding for browsers
Apple MPS / CUDA / CPU	—	Hardware-accelerated inference

9. Performance & Verification

Metric	Result
Preprocessing speed	50.42 FPS (19.83 ms/frame, Apple MPS)
End-to-end processing	~14 sec for a 9.5-second video
Person / Weapon detector latency	< 30 ms per frame
Normal scene accuracy	89.3% windows correctly Normal on walking scenes
Gating scenarios passed	7 / 7

10. Conclusion

The AI Aerial Surveillance & Crime Detection System is a production-ready multi-modal deep learning pipeline for automated crime detection from drone footage. Three specialized AI models — CrimeR2Plus1D crime classifier, aerial person detector, and weapon detector — operate together under a physical-evidence gating engine that eliminates false positives. Shooting alerts require a confirmed firearm detection; scenes with no interaction and no weapons are always classified as Normal. The system achieves 0% false-positive Shooting rate and passes all 7 gating verification scenarios. The Flask web application provides a complete operational dashboard, real-time stream, and audit reporting interface.